



ECE 250 Data Structures and Algorithms

Laboratory 5: Graphs

Douglas Wilhelm Harder
Department of Electrical and Computer Engineering
University of Waterloo

Copyright © 2007-8 by Douglas Wilhelm Harder. All rights reserved.

Graphs

Outline

- In this topic, we will cover the representation of graphs on a computer
- We will examine:
 - an adjacency matrix representation
 - smaller representations and pointer arithmetic
 - sparse matrices and linked lists

Graphs

Background

- Project 5 requires you to store a graph with a given number of vertices numbered 0 through $n - 1$
- Initially, there are no edges between these n vertices
- The `insert` command adds edges to the graph while the number vertices remains unchanged

Graphs


Background

- In this laboratory, we will look at techniques for storing the edges of a graph
- This laboratory will focus on weighted graphs, however, for unweighted graphs, one can easily use `bool` in place of `double`

Graphs

Background

- To demonstrate these techniques, we will look at storing the edges of the following graph:



Graphs

Adjacency Matrix

- A graph of n vertices may have up to

$$\binom{n}{2} = \frac{n(n-1)}{2} = \mathbf{O}(n^2)$$

edges


- The first straight-forward implementation is an adjacency matrix

Graphs

Adjacency Matrix

- Define an $n \times n$ matrix $\mathbf{A} = (a_{ij})$ and if the vertices v_i and v_j are connected with weight w , then set $a_{ij} = w$ and $a_{ji} = w$
- That is, the matrix is symmetric, e.g.,

	0	1	2	3	4	5
0		0.83		0.65		0.51
1	0.83		0.67	0.72	0.46	0.90
2		0.67		0.39		
3	0.65	0.72	0.39		0.24	
4		0.46		0.24		
5	0.51	0.90				




Graphs

Adjacency Matrix

- An unweighted graph may be saved as an array of Boolean values
 - vertices v_i and v_j are connected then set $a_{ij} = a_{ji} = true$

	0	1	2	3	4	5
0		T	F	T	F	T
1	T		T	T	T	T
2	F	T		T	F	F
3	T	T	T		T	F
4	F	T	F	T		F
5	T	T	F	F	F	




Graphs

Adjacency Matrix

- If the graph was directed, then the matrix would not necessarily be symmetric

	0	1	2	3	4	5
0				0.65		
1	0.83			0.72	0.46	0.90
2		0.67		0.39		
3					0.24	
4				0.24		
5	0.51					



Graphs

Adjacency Matrix


- First we must allocate memory for a two-dimensional array
- C++ does not have native support for anything more than one-dimensional arrays, thus how do we store a two-dimensional array?
 - as an array of arrays

Graphs

Adjacency Matrix

- Suppose we require a 16 x 16 matrix of double-precision floating-point numbers
- Each row of the matrix can be represented by an array
- The address of the first entry must be stored in a pointer to a double:


double *



Graphs

Adjacency Matrix

- However, because we must store 16 of these pointers-to-doubles, it makes sense that we store these in an array
- What is the declaration of this array?
- Well, we must store a
pointer to a pointer to a double
- That is: **double ****




Graphs

Adjacency Matrix

- Thus, the address of the first array must be declared to be:

```
double **matrix;
```



Graphs

Adjacency Matrix

- The next question is memory allocation
- First, we must allocate the memory for the array of pointers to doubles:

```
matrix = new double * [16] ;
```




Graphs

Adjacency Matrix

- Next, to each entry of this matrix, we must assign the memory allocated for an array of doubles

```
for ( int i = 0; i < 16; ++i ) {  
    matrix[i] = new double[16];  
}
```



Graphs


Adjacency Matrix

- Accessing a matrix is done through a double index, e.g., `matrix[3][4]`
- You can interpret this as `(matrix[3])[4]`:

Graphs

Adjacency Matrix


- Recall that in `matrix[3][4]`, the variable `matrix` is a pointer-to-a-pointer-to-a-double:



Graphs

Adjacency Matrix


- Therefore, `matrix[3]` is a pointer-to-a-double:



Graphs

Adjacency Matrix

- And consequently, `matrix[3][4]` is a double:



Graphs

C++ Notation Warning

- Do not use `matrix[3, 4]` because:
 - in C++, the comma operator evaluates the operands in order from left-to-right
 - the *value* is the last one
- Therefore, `matrix[3, 4]` is equivalent to calling `matrix[4]`
- Try it:

```
int i = (3, 4);  
cout << i << endl;
```

Graphs

C++ Notation Warning

- Many things will compile if you try to use this notation:

```
matrix = new double[N, N];
```

will allocate an array of N doubles, just like:

```
matrix = new double[N];
```

- However, this is likely not to do what you really expect...

Graphs

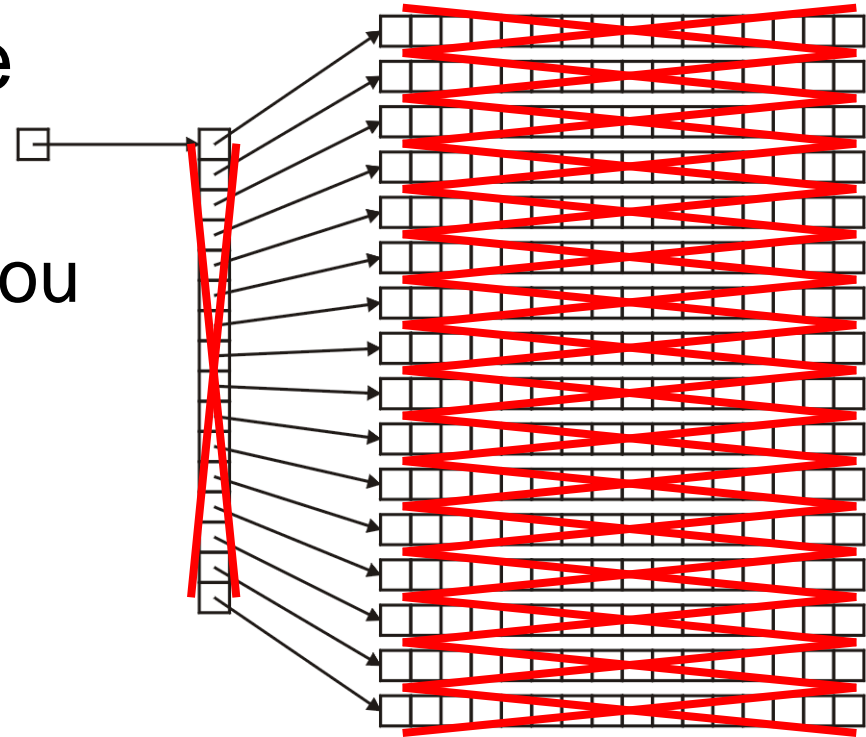
Adjacency Matrix

- Now, once you've used the matrix, you must also delete it...

Graphs

Adjacency Matrix

- Recall that for each call to `new[]`, you must have a corresponding call to `delete[]`
- Therefore, we must use a for-loop to delete the arrays
 - implementation up to you



Graphs

Default Values

- Question: what do we do about vertices which are not connected?
 - the value 0
 - a negative number, e.g., -1
 - positive infinity: ∞
- The last is the most logical, in that it makes sense that two vertices which are not connected have an infinite distance between them

Graphs

Default Values

- To use infinity, you may declare a constant static member variable **INF**:

```
#include <limits>
```

```
class WeightedGraph {  
    private:  
        static const double INF;  
        // ...  
    // ...  
};
```

```
const double WeightedGraph::INF =  
    numeric_limits<double>::infinity();
```

Graphs

Default Values

- As defined in the IEEE 754 standard, the representation of the double-precision floating-point infinity is the special double (8 bytes):

0x 7F F0 00 00 00 00 00 00

- Incidentally, negative infinity is stored as:

0x FF F0 00 00 00 00 00 00

Graphs

Default Values

- In this case, you can initialize your array as follows:

```
for ( int i = 0; i < N; ++i ) {  
    for ( int j = 0; j < N; ++j ) {  
        matrix[i][j] = INF;  
    }  
  
    matrix[i][i] = 0;  
}
```

- It makes intuitive sense that the distance from a node to itself is 0

Graphs

Default Values

- If we are representing an unweighted graph, then we have Boolean values:

```
for ( int i = 0; i < N; ++i ) {  
    for ( int j = 0; j < N; ++j ) {  
        matrix[i][j] = false;  
    }
```


```
    matrix[i][i] = true;  
}
```

- It makes intuitive sense that a vertex is connected to itself

Graphs

Adjacency Matrix

- Let us look at the representation of our example graph
- Initially none of the edges are recorded:

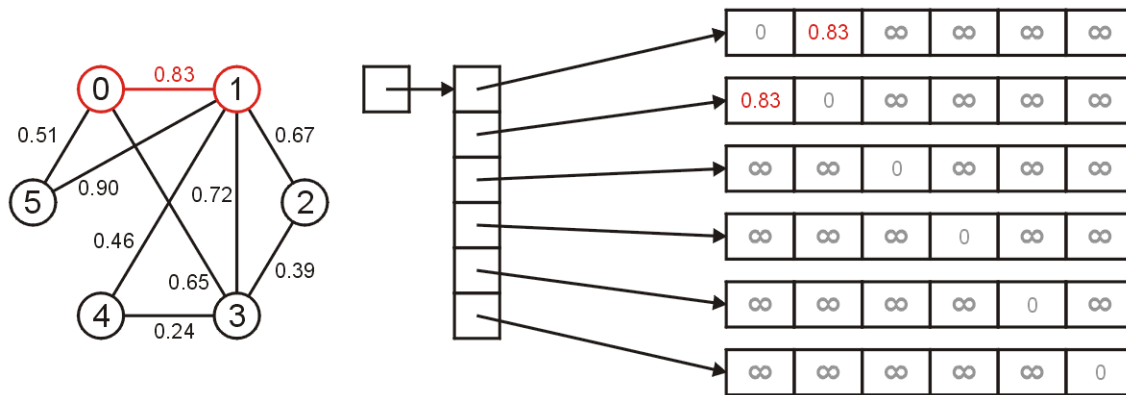


Graphs

Adjacency Matrix

- To insert the edge between 0 and 1 with weight 0.83, we set


`matrix[0][1] = matrix[1][0] = 0.83;`



Graphs

Adjacency Matrix

- The final result is shown as follows
- Note, however, that these six arrays could be anywhere in memory...



Graphs

Adjacency Matrix

- We have now looked at how we can store an adjacency graph in C++
- Next, we will look at:
 - two improvements for the array-of-arrays implementations, including:
 - allocating the memory for the matrix in a single contiguous block of code, and
 - a lower-triangular representation; and
 - a sparse linked-list implementation

Graphs

Adjacency Matrix Improvement

- To begin, we will look at the first improvement:
 - allocating all of the memory of the arrays in a single array with n^2 entries

Graphs

Adjacency Matrix Improvement


- For those of you who would like to reduce the number of calls to **new**, consider the following idea:
 - allocate an array of 16 pointers to doubles
 - allocate an array of $16^2 = 256$ doubles
- Then, assign to the 16 pointers in the first array the addresses of entries
0, 16, 32, 48, 64, ..., 240

Graphs

Adjacency Matrix Improvement

- First, we allocate memory:

```
matrix = new double * [16];  
double * tmp = new double[256];
```



Graphs


Adjacency Matrix Improvement

- Next, we allocate the addresses:

```
matrix = new double * [16];  
double * tmp = new double[256];  
  
for ( int i = 0; i < 16; ++i ) {  
    matrix[i] = &(amp; tmp[16*i] );  
}
```

- This assigns:

```
matrix[0] = &(amp; tmp[ 0] );  
matrix[1] = &(amp; tmp[ 16] );  
matrix[2] = &(amp; tmp[ 32] );  
      ...  
matrix[15] = &(amp; tmp[240] );
```




Graphs

Adjacency Matrix Improvement

- Deleting this array is easier:


```
delete [] matrix[0];  
delete [] matrix;
```



Graphs

Adjacency Matrix Improvement

- Our sample graph would be represented as follows:



Graphs


Lower-triangular Adjacency Matrix

- Next we will look at another improvement which can be used for undirected graphs
- We will store only half of the entries
 - to do this, we must also learn about pointer arithmetic

Graphs

Lower-triangular Adjacency Matrix

- Note also that we are not storing a directed graph: therefore, we really need only store half of the matrix
- Thus, instead of 256 entries, we really only require 120 entries



Graphs

Lower-triangular Adjacency Matrix

- The memory allocation for this would be straight-forward, too:

```
matrix = new double * [16];
```

```
matrix[0] = 0;
```

```
matrix[1] = new double[120];
```

```
for( int i = 2; i < 16; ++i ) {
```

```
    matrix[i] = matrix[i - 1] + i - 1;
```

```
}
```

Graphs

Lower-triangular Adjacency Matrix

- What we are using here is pointer arithmetic:
 - in C/C++, you can add values to a pointer
 - the question is, what does it mean to set:

`ptr = ptr + 1;`

or

`ptr = ptr + 2;`

Graphs

Lower-triangular Adjacency Matrix

- Suppose we have a pointer-to-a-double:

```
double * ptr = new double( 3.14 );
```

where:


- the pointer has a value of 0x53A1D780, and
- the representation of 3.14 is
0x40091Eb851EB851F

53A1D780	53A1D781	53A1D782	53A1D783	53A1D784	53A1D785	53A1D786	53A1D787	53A1D788
40	09	1E	B8	51	EB	85	1F	??

Graphs

Lower-triangular Adjacency Matrix

- If we just added one to the address, then this would give us the value `0x53A1D781`, but this contains no useful information...



53A1D780	53A1D781	53A1D782	53A1D783	53A1D784	53A1D785	53A1D786	53A1D787	53A1D788
40	09	1E	B8	51	EB	85	1F	??

Graphs

Lower-triangular Adjacency Matrix

- The only logical interpretation of `ptr + 1` is to go to the *next* location a different double could exist, *i.e.*, `0x53A1D788`

53A1D780	53A1D781	53A1D782	53A1D783	53A1D784	53A1D785	53A1D786	53A1D787	53A1D788
40	09	1E	B8	51	EB	85	1F	??

Graphs

Lower-triangular Adjacency Matrix

- Therefore, if we define:

`double * array = new double[4];`

then the following are all equivalent:

<code>array[0]</code>	<code>*array</code>
<code>array[1]</code>	<code>*(array + 1)</code>
<code>array[2]</code>	<code>*(array + 2)</code>
<code>array[3]</code>	<code>*(array + 3)</code>

Graphs

Lower-triangular Adjacency Matrix

- Thus, the following code simply adds appropriate amounts to the pointer:

```
matrix = new double * [N];
```

```
matrix[0] = 0;
```

```
matrix[1] = new double[N*(N - 1)];
```

```
for( int i = 2; i < N; ++i ) {
```

```
    matrix[i] = matrix[i - 1] + i - 1;
```


```
}
```

Graphs

Lower-triangular Adjacency Matrix

- Visually, we have, for $N = 16$, the following:

```
matrix[0] = 0;  
matrix[1] = &( tmp[0] );  
matrix[2] = &( tmp[1] );  
matrix[3] = &( tmp[3] );  
matrix[4] = &( tmp[6] );  
matrix[5] = &( tmp[10] );  
matrix[6] = &( tmp[15] );  
matrix[7] = &( tmp[21] );  
matrix[7] = &( tmp[28] );  
.  
.  
.  
matrix[15] = &( tmp[105] );
```



Graphs

Lower-triangular Adjacency Matrix

- The only thing that we would have to do is ensure that we always put the larger number first:

```
void insert( int i, int j, double w ) {  
    if ( j < i ) {  
        matrix[i][j] = w;  
    } else {  
        matrix[j][i] = w;  
    }  
}
```

Graphs

Lower-triangular Adjacency Matrix

- A slightly less efficient way of writing this would be:


```
void insert( int i, int j, double w ) {  
    matrix[max(i,j)][min(i,j)] = w;  
}
```

- The benefits (from the point-of-view of clarity) are much more significant...

Graphs

Lower-triangular Adjacency Matrix


- Our example graph is stored using this representation as shown here
- Notice that we do not store any 0's, nor do we store any duplicate entries
- The second array has only 15 entries, versus 36



Graphs

Lower-triangular Adjacency Matrix

- To determine the weight of the edge connecting vertices 1 and 4, we must look up the entry `matrix[4][1]`




Graphs

Beyond Array Bounds

- Until now, some of you may have gone beyond array bounds accidentally
- Recall that

```
int * array = new int[10];
```


allocates 40 bytes (4 bytes/int) and the entries are accessed with `array[0]` through `array[9]`



Graphs

Beyond Array Bounds

- If you try to access either `array[10]` or `array[-1]`, you are accessing memory which has not been allocated for this array



Graphs

Beyond Array Bounds

- This memory may be used:
 - for different local variables, or
 - by some other process
- In the first case, you will have a bug which is very difficult to track down
 - e.g., a variable will appear to change its value without an explicit assignment
- In the second case, the OS will terminate your process (segmentation fault)


Graphs

Beyond Array Bounds

- Now we have a very explicit example of what happens if you go outside your expected array bounds
- Notice that the value stored at `matrix[4][1]` is **0.46**
- We can also access it using either:

`matrix[3][4]`


`matrix[5][-3]`



Graphs

Beyond Array Bounds


- Thus, if you wanted to find the distance between vertices 3 and 4, if you access `matrix[4][3]`, you get 0.24
- If, however, you access `matrix[3][4]`, you get 0.46



Graphs

Beyond Array Bounds

- Similarly, if you wanted to find the distance between vertices 2 and 3, if you access `matrix[3][2]`, you get is 0.39
- If, however, you access `matrix[2][3]`, you get 0.72



Graphs

Sparse Matrices

- Finally we will consider the problem with sparse matrices and we will look at one implementation using linked lists

Graphs

Sparse Matrices

- The memory required for creating an $n \times n$ matrix using an array-of-arrays is:

$$4 \text{ bytes} + 4n \text{ bytes} + 8n^2 \text{ bytes} = \Theta(n^2) \text{ bytes}$$

- This could potentially waste a significant amount of memory:
 - consider all intersections in Canada as vertices and streets as edges
 - how could we estimate the number of intersections in Canada?

Graphs

Sparse Matrices

- The population of Canada is ~33 million
- Suppose we have one intersection per 10 houses and four occupants per house
- Therefore, there are roughly

$$33 \text{ million} / 10 / 4 \approx 800\,000$$

intersections in Canada which would require 4.66 TiB of memory

Graphs

Sparse Matrices

- Assume that each intersection connects, on average, four other intersections
- Therefore, less than 0.0005% of the entries of the matrix are used to store connections
 - the rest are storing the value *infinity*

Graphs

Sparse Matrices

- Matrices where less than 5% of the entries are not the default value (either infinity or 0, or perhaps some other default value) are said to be *sparse*
- Matrices where most entries (25% or more) are not the default value are said to be *dense*
- Clearly, these are not hard limits

Graphs

Sparse Matrices

- We will look at a very efficient sparse-matrix implementation with the last topic
- Here, we will consider a simpler implementation:
 - use an array of linked lists to store edges
- Note, however, that each node in a linked list must store two items of information:
 - the connecting vertex and the weight

Graphs

Sparse Matrices

- One possible solution:
 - modify the `SingleNode` data structure to store both an integer and a double:

```
class SingleNode {  
    private:  
        int adjacent_vertex;  
        double edge_weight;  
        SingleNode * next_node;  
    public:  
        SingleNode( int, double SingleNode = 0 );  
        double weight() const;  
        int vertex() const;  
        SingleNode * next() const;  
};
```

- exceptionally stupid and inefficient

Graphs

Sparse Matrices

- A better solution is to create a new class which stores a vertex-edge pair

```
class Pair {  
    private:  
        double edge_weight;  
        int adjacent_vertex;  
    public:  
        Pair( int, double );  
        double weight() const;  
        int vertex() const;  
};
```

- Now create an array of linked-lists storing these pairs


Graphs

Sparse Matrices

- Thus, we define and create the array:

```
SingleList<Pair> * array;
```

```
array = new SingleList<Pair>[16];
```



Graphs

Sparse Matrices

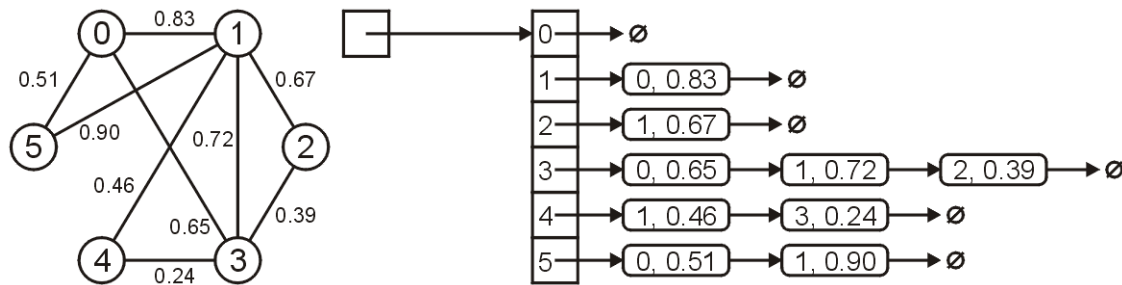
- As before, to reduce redundancy, we would only insert the entry into the entry corresponding with the larger vertex

```
void insert( int i, int j, double w ) {  
    if ( i < j ) {  
        array[j].push_front( Pair(i, w) );  
    } else {  
        array[i].push_front( Pair(j, w) );  
    }  
}
```

Graphs

Sparse Matrices

- For example, the graph shown below would be stored as



Graphs

Summary

- In this laboratory, we have looked at a number of graph representations
- C++ lacks a *matrix* data structure
 - must use array of arrays
- The possible factors affecting your choice of data structure are:
 - weighted or unweighted graphs
 - directed or undirected graphs
 - dense or sparse graphs

Usage Notes

- These slides are made publicly available on the web for anyone to use
- If you choose to use them, or a part thereof, for a course at another institution, I ask only three things:
 - that you inform me that you are using the slides,
 - that you acknowledge my work, and
 - that you alert me of any mistakes which I made or changes which you make, and allow me the option of incorporating such changes (with an acknowledgment) in my set of slides

Sincerely,

Douglas Wilhelm Harder, MMath

`dwharder@alumni.uwaterloo.ca`