

Pre-lab #2 tutorial

ECE 254
Operating Systems and Systems Programming

May 24, 2012

Content

- Concurrency
- Concurrent Programming
- Thread vs. Process
- POSIX Threads
- Synchronization and Critical Sections
- Mutexes
- Semaphores

Concurrency

- Concurrency is a property where multiple paths of execution are running simultaneously
- Multiple processes/threads running on a single processor
- Multiple processes/threads running on multiple processors/machines (true concurrency)
- Requires a different programming paradigm: Concurrent programming

Concurrent Programming

- A programming paradigm that involves designing programs so that they can run in parallel
- Concurrent programs can share data and resources
- Two mechanisms of communication among concurrent components:
 - Message passing
 - *Shared memory*

Concurrent Programming (Cont.)

- Main challenges of concurrent programming:
 - Ensuring correct and meaningful order of execution
 - Coordinating access to shared data and resources to ensure their consistency throughout execution

Thread vs. Process

- A thread is a mechanism to allow a program to run things in parallel (like processes)
- A thread is a finer-grained unit of execution than a process. Each running process has at least one thread
- Each process has its own memory space, variables etc.
- Threads in a process share the same memory space, variables etc.

POSIX Threads

- POSIX (**P**ortable **O**perating **S**ystem **I**nterface) is a group of standards for compatability among different operating systems
- POSIX Threads (*pthread*s) is a standard for creating and using threads
- We will use the GNU/Linux implementation of the *pthread*s

Thread Creation

- We will mainly use three functions: *pthread_create* , *pthread_join* and *pthread_cancel*
- We use the *pthread_create* to create a new thread, it takes four parameters as follows:
 - A pointer to *pthread_t* variable which stores the thread ID
 - A pointer to a thread attribute object which specifies how the thread interacts with the program, or a *NULL* value
 - Pointer to the function the thread will execute. The function must have a return type *void** and takes *void** as a parameter
 - A thread argument value of type *void** or *NULL* value

Thread Creation (Cont.)

```
#include <pthread.h>
#include <stdio.h>

void*
print_func (void* args)
{
    int count;
    int i;

    count = *(int *) args;
    for (i = 0; i < count; ++i)
        fprintf (stdout, "Hello world!\n");

    return NULL;
}

int
main (int argc, char **argv)
{
    pthread_t thread_id;
    int count;
    int i;

    count = 10000;

    pthread_create (&thread_id, NULL, &print_func, (void*) &count);
    pthread_join (thread_id, NULL);

    return 0;
}
```

Thread Creation (Cont.)

- We use the *pthread_join* to ensure that the thread executing *main* will wait for the other thread to finish
- This ensures that the created thread will no be using deallocated variables

Thread Cancellation

- One thread can request the termination of another thread by calling *pthread_cancel*, passing the thread ID as an argument
- A thread can have one of three different “cancellation” states:
 - Asynchronously cancelable
 - Synchronously cancelable
 - Uncancelable
- A thread can alter its “cancellation” at any time using the *pthread_setcanceltype* function

Synchronization and Critical Sections

- There is no way to tell when a specific thread will be scheduled or at which line of code the OS will suspend one thread and re-schedule another
- If there is a concurrency bug, it is very hard to fix since it won't be deterministic
- The main cause of concurrency bugs is threads attempting to access the same resource/variable at the same time. This is called a *race condition*
- Race conditions can crash a program, or even worse leave variables in a corrupted state which could radically alter the behavior of the program

Race Condition

```
#include <malloc.h>

struct job
{
    struct job *next;
};

struct job *job_queue;

void*
thread_function (void *arg)
{
    while (job_queue != NULL)
    {
        struct job *next_job = job_queue;

        job_queue = job_queue->next;
        process_job (next_job);

        free (next_job);
    }

    return NULL;
}
```

Atomicity

- To avoid concurrency bugs and race conditions, you have to make sure that operations on shared variables are *atomic*
- An atomic operation means it is totally indivisible and uninterruptible
- To ensure atomicity we will use:
 - Mutexes
 - Semaphores

Atomicity (Cont.)

- Mutex stands for *MUTual EXclusion locks*
- A special lock that one thread can lock at a time, before entering a *critical section*
- Any thread attempting to lock an already locked mutex will be blocked
- The underlying OS ensures that no race conditions will occur among threads using the mutex
- To create a mutex on Linux:

```
pthread_mutex_t mutex;  
pthread_mutex_init (&mutex, NULL);
```

- To lock/unlock a mutex:

```
pthread_mutex_lock (&mutex)  
pthread_mutex_unlock (&mutex)
```

Mutex in action

```
#include <malloc.h>
#include <pthread.h>

.....
pthread_mutex_t job_queue_mutex;
pthread_mutex_init (&job_queue_mutex, NULL);

void*
thread_function (void *arg)
{
    while (1)
    {
        struct job* next_job;

        /** Start of critical section **/
        pthread_mutex_lock (&job_queue_mutex);
        if (job_queue == NULL)
        {
            next_job = NULL;
        }
        else
        {
            next_job = job_queue;
            job_queue = job_queue->next;
        }
        pthread_mutex_unlock (&job_queue_mutex);
        /** End of critical section **/

        process_job (next_job);
        free (next_job);
    }
    return NULL;
}
```


Deadlocks

- Mutexes introduce another type of concurrency bugs: *deadlocks*
- One or more threads can block forever waiting for a mutex to be unlocked
- A deadlock can easily take place if a thread who originally locked a mutex tries to lock it again
- In Linux, this can happen with the default type of mutexes, the *fast mutex*
- The *recursive mutex*, on the other hand, can be locked any number of times, but it must also be unlocked the same number of times

Semaphores

- In the previous queue example, what if all the jobs are not queued in advance?
- There is a possibility that threads can finish processing before all jobs are pushed onto the queue
- We need a *semaphore* to synchronize the threads, so that they can “wait” if there are no jobs on the queue at the moment

Semaphores (Cont.)

- A semaphore stores a non-negative counter which can be incremented/decremented. It supports two main operations:
 - *wait*: decrements the value of a semaphore by one. If the value is already zero, the function blocks and hence the calling thread
 - *post*: increments the value of a semaphore by one. If the value was originally zero, one of the blocked threads (if any) will be allowed to run
- We will use the POSIX standard semaphore for thread synchronization on Linux. There is another type of Linux semaphore used with processes (not covered)
- Semaphore common operations:

```
sem_wait
sem_post
sem_trywait    // Non blocking wait
sem_getvalue   // Get value of the semaphore
```

Semaphores Example

```
#include <malloc.h>
#include <pthread.h>
#include <semaphore.h>

struct job
{
    struct job *next;
};

struct job *job_queue;

pthread_mutex_t job_queue_mutex;
pthread_mutex_init (&job_queue_mutex, NULL);

sem_t job_queue_count;

void
initialize_job_queue ()
{
    job_queue = NULL;
    sem_init (&job_queue_count, 0, 0);
}

.....
```

Semaphores Example (cont.)

```
.....  
  
void*  
thread_function (void *arg)  
{  
    while (1)  
    {  
        struct job *next_job;  
  
        sem_wait (&job_queue_count);  
  
        /** Start of critical section **/  
        pthread_mutex_lock (&job_queue_mutex);  
        next_job = job_queue;  
        job_queue = job_queue->next;  
        pthread_mutex_unlock (&job_queue_mutex);  
        /** End of critical section **/  
  
        process_job (next_job);  
        free (next_job);  
    }  
    return NULL;  
}  
  
.....
```

Semaphores Example (cont.)

.....

```
void
enqueue_job (/* Pass job-specific data here */)
{
    struct job *new_job;

    new_job = (struct job*) malloc (sizeof (struct job));

    /** Start of critical section **/
    pthread_mutex_lock (&job_queue_mutex);
    new_job->next = job_queue;
    job_queue = new_job;

    sem_post (&job_queue_count);

    pthread_mutex_unlock (&job_queue_mutex);
    /** End of critical section **/
}
```

Questions

Questions ?

References

- M. Mitchell, J. Oldham, and A. Samuel (2001). *Threads , Advanced Linux Programming*. Indianapolis:New Riders Publishing